

## Problem Statement:

A banking institution requires actionable insights from the perspective of Mortgage-Backed Securities, Geographic Business Investment and Real Estate Analysis.

The objective is to identify white spaces/potential business in the mortgage loan. The mortgage bank would like to identify potential monthly mortgage expenses for each of region based on factors which are primarily monthly family income in a region and rented value of the real estate. Some of the regions are growing rapidly and Competitor banks are selling mortgage loans to subprime customers at a lower interest rate. The bank is strategizing for better market penetration and targeting new customers. A statistical model needs to be created to predict the potential demand in dollars amount of loan for each of the region in the USA. Also, there is a need to create a dashboard which would refresh periodically post data retrieval from the agencies. This would help to monitor the key metrics and trends.

The dashboard must demonstrate relationships and trends for the key metrics as follows: number of loans, average rental income, monthly mortgage and owner's cost, family income vs mortgage cost comparison across different regions. The metrics are described not to limit the dashboard to these few only.

## Dataset Description

Following are the themes the fields fall under Home Owner Costs: Sum of utilities, property taxes.

Second Mortgage: Households with a second mortgage statistics.

Home Equity Loan: Households with a Home equity Loan statistics.

Debt: Households with any type of debt statistics.

Mortgage Costs: Statistics regarding mortgage payments, home equity loans, utilities and property taxes

Home Owner Costs: Sum of utilities, property taxes statistics

Gross Rent: Contract rent plus the estimated average monthly cost of utility features

Gross Rent as Percent of Income Gross rent as the percent of income very interesting

High school Graduation: High school graduation statistics.

Population Demographics: Population demographic statistics.

Age Demographics: Age demographic statistics.

Household Income: Total income of people residing in the household.

Family Income: Total income of people related to the householder

## Project Task: Week 1

### Import Required Libraries

In [1]:

```
import time
import random
from math import *
import operator
import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
%matplotlib inline
from pandas.plotting import scatter_matrix
import seaborn as sns

import warnings
warnings.filterwarnings("ignore")
```

## Data Import and Preparation:

Import data.

```
In [2]: df_train = pd.read_csv("train.csv")
```

```
In [3]: df_test = pd.read_csv("test.csv")
```

```
In [4]: df_train.head()
```

```
Out[4]:
```

|   | UID    | BLOCKID | SUMLEVEL | COUNTYID | STATEID | state       | state_ab | city       | p     |
|---|--------|---------|----------|----------|---------|-------------|----------|------------|-------|
| 0 | 267822 | NaN     | 140      | 53       | 36      | New York    | NY       | Hamilton   | Ham   |
| 1 | 246444 | NaN     | 140      | 141      | 18      | Indiana     | IN       | South Bend | Rose  |
| 2 | 245683 | NaN     | 140      | 63       | 18      | Indiana     | IN       | Danville   | Dai   |
| 3 | 279653 | NaN     | 140      | 127      | 72      | Puerto Rico | PR       | San Juan   | Guay  |
| 4 | 247218 | NaN     | 140      | 161      | 20      | Kansas      | KS       | Manhattan  | Manh. |

5 rows × 80 columns

```
In [5]: df_train.shape
```

```
Out[5]: (27321, 80)
```

```
In [6]: df_test.head()
```

Out [6]:

|   | UID    | BLOCKID | SUMLEVEL | COUNTYID | STATEID | state        | state_ab | city           |
|---|--------|---------|----------|----------|---------|--------------|----------|----------------|
| 0 | 255504 | NaN     | 140      | 163      | 26      | Michigan     | MI       | Detroit        |
| 1 | 252676 | NaN     | 140      | 1        | 23      | Maine        | ME       | Auburn         |
| 2 | 276314 | NaN     | 140      | 15       | 42      | Pennsylvania | PA       | Pine City      |
| 3 | 248614 | NaN     | 140      | 231      | 21      | Kentucky     | KY       | Monticello     |
| 4 | 286865 | NaN     | 140      | 355      | 48      | Texas        | TX       | Corpus Christi |

5 rows x 80 columns

In [7]: `df_test.shape`

Out[7]: (11709, 80)

In [8]: `df_train.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 27321 entries, 0 to 27320
Data columns (total 80 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   UID                                   27321 non-null  int64
1   BLOCKID                             0 non-null      float64
2   SUMLEVEL                             27321 non-null  int64
3   COUNTYID                             27321 non-null  int64
4   STATEID                              27321 non-null  int64
5   state                                27321 non-null  object
6   state_ab                             27321 non-null  object
7   city                                 27321 non-null  object
8   place                                27321 non-null  object
9   type                                  27321 non-null  object
10  primary                              27321 non-null  object
11  zip_code                             27321 non-null  int64
12  area_code                             27321 non-null  int64
13  lat                                   27321 non-null  float64
14  lng                                   27321 non-null  float64
15  ALand                                27321 non-null  float64
16  AWater                                27321 non-null  int64
17  pop                                   27321 non-null  int64
18  male_pop                             27321 non-null  int64
19  female_pop                           27321 non-null  int64
20  rent_mean                             27007 non-null  float64
21  rent_median                           27007 non-null  float64
22  rent_stdev                             27007 non-null  float64
23  rent_sample_weight                     27007 non-null  float64
24  rent_samples                           27007 non-null  float64
```

|    |                             |       |          |         |
|----|-----------------------------|-------|----------|---------|
| 25 | rent_gt_10                  | 27007 | non-null | float64 |
| 26 | rent_gt_15                  | 27007 | non-null | float64 |
| 27 | rent_gt_20                  | 27007 | non-null | float64 |
| 28 | rent_gt_25                  | 27007 | non-null | float64 |
| 29 | rent_gt_30                  | 27007 | non-null | float64 |
| 30 | rent_gt_35                  | 27007 | non-null | float64 |
| 31 | rent_gt_40                  | 27007 | non-null | float64 |
| 32 | rent_gt_50                  | 27007 | non-null | float64 |
| 33 | universe_samples            | 27321 | non-null | int64   |
| 34 | used_samples                | 27321 | non-null | int64   |
| 35 | hi_mean                     | 27053 | non-null | float64 |
| 36 | hi_median                   | 27053 | non-null | float64 |
| 37 | hi_stdev                    | 27053 | non-null | float64 |
| 38 | hi_sample_weight            | 27053 | non-null | float64 |
| 39 | hi_samples                  | 27053 | non-null | float64 |
| 40 | family_mean                 | 27023 | non-null | float64 |
| 41 | family_median               | 27023 | non-null | float64 |
| 42 | family_stdev                | 27023 | non-null | float64 |
| 43 | family_sample_weight        | 27023 | non-null | float64 |
| 44 | family_samples              | 27023 | non-null | float64 |
| 45 | hc_mortgage_mean            | 26748 | non-null | float64 |
| 46 | hc_mortgage_median          | 26748 | non-null | float64 |
| 47 | hc_mortgage_stdev           | 26748 | non-null | float64 |
| 48 | hc_mortgage_sample_weight   | 26748 | non-null | float64 |
| 49 | hc_mortgage_samples         | 26748 | non-null | float64 |
| 50 | hc_mean                     | 26721 | non-null | float64 |
| 51 | hc_median                   | 26721 | non-null | float64 |
| 52 | hc_stdev                    | 26721 | non-null | float64 |
| 53 | hc_samples                  | 26721 | non-null | float64 |
| 54 | hc_sample_weight            | 26721 | non-null | float64 |
| 55 | home_equity_second_mortgage | 26864 | non-null | float64 |
| 56 | second_mortgage             | 26864 | non-null | float64 |
| 57 | home_equity                 | 26864 | non-null | float64 |
| 58 | debt                        | 26864 | non-null | float64 |
| 59 | second_mortgage_cdf         | 26864 | non-null | float64 |
| 60 | home_equity_cdf             | 26864 | non-null | float64 |
| 61 | debt_cdf                    | 26864 | non-null | float64 |
| 62 | hs_degree                   | 27131 | non-null | float64 |
| 63 | hs_degree_male              | 27121 | non-null | float64 |
| 64 | hs_degree_female            | 27098 | non-null | float64 |
| 65 | male_age_mean               | 27132 | non-null | float64 |
| 66 | male_age_median             | 27132 | non-null | float64 |
| 67 | male_age_stdev              | 27132 | non-null | float64 |
| 68 | male_age_sample_weight      | 27132 | non-null | float64 |
| 69 | male_age_samples            | 27132 | non-null | float64 |
| 70 | female_age_mean             | 27115 | non-null | float64 |
| 71 | female_age_median           | 27115 | non-null | float64 |
| 72 | female_age_stdev            | 27115 | non-null | float64 |
| 73 | female_age_sample_weight    | 27115 | non-null | float64 |
| 74 | female_age_samples          | 27115 | non-null | float64 |
| 75 | pct_own                     | 27053 | non-null | float64 |
| 76 | married                     | 27130 | non-null | float64 |
| 77 | married_snp                 | 27130 | non-null | float64 |
| 78 | separated                   | 27130 | non-null | float64 |
| 79 | divorced                    | 27130 | non-null | float64 |

dtypes: float64(62), int64(12), object(6)

memory usage: 16.7+ MB

In [9]: `df_test.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 11709 entries, 0 to 11708
Data columns (total 80 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   UID                                    11709 non-null  int64
1   BLOCKID                               0 non-null      float64
2   SUMLEVEL                              11709 non-null  int64
3   COUNTYID                              11709 non-null  int64
4   STATEID                               11709 non-null  int64
5   state                                 11709 non-null  object
6   state_ab                              11709 non-null  object
7   city                                  11709 non-null  object
8   place                                 11709 non-null  object
9   type                                  11709 non-null  object
10  primary                               11709 non-null  object
11  zip_code                              11709 non-null  int64
12  area_code                             11709 non-null  int64
13  lat                                    11709 non-null  float64
14  lng                                    11709 non-null  float64
15  ALand                                  11709 non-null  int64
16  AWater                                11709 non-null  int64
17  pop                                    11709 non-null  int64
18  male_pop                              11709 non-null  int64
19  female_pop                            11709 non-null  int64
20  rent_mean                             11561 non-null  float64
21  rent_median                           11561 non-null  float64
22  rent_stdev                             11561 non-null  float64
23  rent_sample_weight                    11561 non-null  float64
24  rent_samples                           11561 non-null  float64
25  rent_gt_10                             11560 non-null  float64
26  rent_gt_15                             11560 non-null  float64
27  rent_gt_20                             11560 non-null  float64
28  rent_gt_25                             11560 non-null  float64
29  rent_gt_30                             11560 non-null  float64
30  rent_gt_35                             11560 non-null  float64
31  rent_gt_40                             11560 non-null  float64
32  rent_gt_50                             11560 non-null  float64
33  universe_samples                       11709 non-null  int64
34  used_samples                           11709 non-null  int64
35  hi_mean                                11587 non-null  float64
36  hi_median                              11587 non-null  float64
37  hi_stdev                               11587 non-null  float64
38  hi_sample_weight                       11587 non-null  float64
39  hi_samples                             11587 non-null  float64
40  family_mean                            11573 non-null  float64
41  family_median                          11573 non-null  float64
42  family_stdev                           11573 non-null  float64
43  family_sample_weight                    11573 non-null  float64
44  family_samples                          11573 non-null  float64
45  hc_mortgage_mean                       11441 non-null  float64
46  hc_mortgage_median                     11441 non-null  float64
47  hc_mortgage_stdev                       11441 non-null  float64
```

```

48 hc_mortgage_sample_weight    11441 non-null float64
49 hc_mortgage_samples          11441 non-null float64
50 hc_mean                      11419 non-null float64
51 hc_median                    11419 non-null float64
52 hc_stdev                     11419 non-null float64
53 hc_samples                   11419 non-null float64
54 hc_sample_weight             11419 non-null float64
55 home_equity_second_mortgage  11489 non-null float64
56 second_mortgage              11489 non-null float64
57 home_equity                  11489 non-null float64
58 debt                         11489 non-null float64
59 second_mortgage_cdf          11489 non-null float64
60 home_equity_cdf              11489 non-null float64
61 debt_cdf                     11489 non-null float64
62 hs_degree                    11624 non-null float64
63 hs_degree_male               11620 non-null float64
64 hs_degree_female             11604 non-null float64
65 male_age_mean                11625 non-null float64
66 male_age_median              11625 non-null float64
67 male_age_stdev               11625 non-null float64
68 male_age_sample_weight       11625 non-null float64
69 male_age_samples             11625 non-null float64
70 female_age_mean              11613 non-null float64
71 female_age_median            11613 non-null float64
72 female_age_stdev             11613 non-null float64
73 female_age_sample_weight     11613 non-null float64
74 female_age_samples           11613 non-null float64
75 pct_own                      11587 non-null float64
76 married                      11625 non-null float64
77 married_snp                  11625 non-null float64
78 separated                    11625 non-null float64
79 divorced                     11625 non-null float64
dtypes: float64(61), int64(13), object(6)
memory usage: 7.1+ MB

```

Figure out the primary key and look for the requirement of indexing.

```

In [10]: #UID is the primary key

df_train.set_index(keys=['UID'], inplace=True)

df_test.set_index(keys=['UID'], inplace=True)

```

```

In [11]: df_train.head(2)

```

Out[11]:

|  | BLOCKID | SUMLEVEL | COUNTYID | STATEID | state | state_ab | city | place      |
|--|---------|----------|----------|---------|-------|----------|------|------------|
|  | UID     |          |          |         |       |          |      |            |
|  | 267822  | NaN      | 140      | 53      | 36    | New York | NY   | Hamilton   |
|  | 246444  | NaN      | 140      | 141     | 18    | Indiana  | IN   | South Bend |

2 rows x 79 columns

In [12]:

```
df_test.head(2)
```

Out[12]:

|  | BLOCKID | SUMLEVEL | COUNTYID | STATEID | state | state_ab | city | place   |
|--|---------|----------|----------|---------|-------|----------|------|---------|
|  | UID     |          |          |         |       |          |      |         |
|  | 255504  | NaN      | 140      | 163     | 26    | Michigan | MI   | Detroit |
|  | 252676  | NaN      | 140      | 1       | 23    | Maine    | ME   | Auburn  |

2 rows x 79 columns

Gauge the fill rate of the variables and devise plans for missing value treatment. Please explain explicitly the reason for the treatment chosen for each variable.

In [13]:

```
# percentage of missing values in train set
missing_list_train = df_train.isnull().sum()*100/len(df_train)

missing_values_df_train = pd.DataFrame(missing_list_train, columns=['Percentage_of_missing_values'])

missing_values_df_train.sort_values(by = ['Percentage_of_missing_values'],
missing_values_df_train[missing_values_df_train['Percentage_of_missing_val
```

Out[13]:

|                  | Percentage_of_missing_values |
|------------------|------------------------------|
| BLOCKID          | 100.000000                   |
| hc_samples       | 2.196113                     |
| hc_mean          | 2.196113                     |
| hc_median        | 2.196113                     |
| hc_stddev        | 2.196113                     |
| hc_sample_weight | 2.196113                     |
| hc_mortgage_mean | 2.097288                     |



|                             |          |
|-----------------------------|----------|
| hc_mortgage_stdev           | 2.097288 |
| hc_mortgage_sample_weight   | 2.097288 |
| hc_mortgage_samples         | 2.097288 |
| hc_mortgage_median          | 2.097288 |
| home_equity_second_mortgage | 1.672706 |
| home_equity                 | 1.672706 |
| debt                        | 1.672706 |
| second_mortgage_cdf         | 1.672706 |
| home_equity_cdf             | 1.672706 |
| debt_cdf                    | 1.672706 |
| second_mortgage             | 1.672706 |
| rent_gt_15                  | 1.149299 |
| rent_gt_50                  | 1.149299 |
| rent_gt_40                  | 1.149299 |
| rent_gt_35                  | 1.149299 |
| rent_gt_30                  | 1.149299 |
| rent_gt_25                  | 1.149299 |
| rent_gt_20                  | 1.149299 |
| rent_samples                | 1.149299 |
| rent_gt_10                  | 1.149299 |
| rent_sample_weight          | 1.149299 |
| rent_stdev                  | 1.149299 |
| rent_median                 | 1.149299 |
| rent_mean                   | 1.149299 |
| family_median               | 1.090736 |
| family_samples              | 1.090736 |
| family_sample_weight        | 1.090736 |
| family_stdev                | 1.090736 |
| family_mean                 | 1.090736 |
| hi_stdev                    | 0.980930 |
| hi_sample_weight            | 0.980930 |
| hi_samples                  | 0.980930 |
| hi_median                   | 0.980930 |
| hi_mean                     | 0.980930 |

|                          |          |
|--------------------------|----------|
| pct_own                  | 0.980930 |
| hs_degree_female         | 0.816222 |
| female_age_samples       | 0.753999 |
| female_age_sample_weight | 0.753999 |
| female_age_stdev         | 0.753999 |
| female_age_median        | 0.753999 |
| female_age_mean          | 0.753999 |
| hs_degree_male           | 0.732038 |
| separated                | 0.699096 |
| married_snp              | 0.699096 |
| married                  | 0.699096 |
| divorced                 | 0.699096 |
| hs_degree                | 0.695436 |
| male_age_stdev           | 0.691776 |
| male_age_samples         | 0.691776 |
| male_age_mean            | 0.691776 |
| male_age_median          | 0.691776 |
| male_age_sample_weight   | 0.691776 |

In [14]:

```
# percentage of missing values in test set
missing_list_test = df_test.isnull().sum()*100/len(df_train)

missing_values_df_test = pd.DataFrame(missing_list_test, columns=['Percentage_of_missing_values'])

missing_values_df_test.sort_values(by = ['Percentage_of_missing_values'], ascending=False)

missing_values_df_test[missing_values_df_test['Percentage_of_missing_values'] > 0]
```

Out[14]:

|                   | Percentage_of_missing_values |
|-------------------|------------------------------|
| BLOCKID           | 42.857143                    |
| hc_samples        | 1.061455                     |
| hc_mean           | 1.061455                     |
| hc_median         | 1.061455                     |
| hc_stdev          | 1.061455                     |
| hc_sample_weight  | 1.061455                     |
| hc_mortgage_mean  | 0.980930                     |
| hc_mortgage_stdev | 0.980930                     |

|                             |          |
|-----------------------------|----------|
| hc_mortgage_sample_weight   | 0.980930 |
| hc_mortgage_samples         | 0.980930 |
| hc_mortgage_median          | 0.980930 |
| home_equity_second_mortgage | 0.805241 |
| home_equity                 | 0.805241 |
| debt                        | 0.805241 |
| second_mortgage_cdf         | 0.805241 |
| home_equity_cdf             | 0.805241 |
| debt_cdf                    | 0.805241 |
| second_mortgage             | 0.805241 |
| rent_gt_20                  | 0.545368 |
| rent_gt_50                  | 0.545368 |
| rent_gt_40                  | 0.545368 |
| rent_gt_35                  | 0.545368 |
| rent_gt_30                  | 0.545368 |
| rent_gt_25                  | 0.545368 |
| rent_gt_10                  | 0.545368 |
| rent_gt_15                  | 0.545368 |
| rent_samples                | 0.541708 |
| rent_sample_weight          | 0.541708 |
| rent_stddev                 | 0.541708 |
| rent_median                 | 0.541708 |
| rent_mean                   | 0.541708 |
| family_median               | 0.497786 |
| family_samples              | 0.497786 |
| family_sample_weight        | 0.497786 |
| family_stddev               | 0.497786 |
| family_mean                 | 0.497786 |
| hi_stddev                   | 0.446543 |
| hi_median                   | 0.446543 |
| pct_own                     | 0.446543 |
| hi_mean                     | 0.446543 |
| hi_sample_weight            | 0.446543 |
| hi_samples                  | 0.446543 |

|                          |          |
|--------------------------|----------|
| hs_degree_female         | 0.384320 |
| female_age_mean          | 0.351378 |
| female_age_sample_weight | 0.351378 |
| female_age_median        | 0.351378 |
| female_age_stdev         | 0.351378 |
| female_age_samples       | 0.351378 |
| hs_degree_male           | 0.325757 |
| hs_degree                | 0.311116 |
| married_snp              | 0.307456 |
| male_age_samples         | 0.307456 |
| male_age_sample_weight   | 0.307456 |
| male_age_stdev           | 0.307456 |
| male_age_median          | 0.307456 |
| male_age_mean            | 0.307456 |
| married                  | 0.307456 |
| separated                | 0.307456 |
| divorced                 | 0.307456 |

BLOCKID has 100% missing values in train set, and 42% in test set we can drop it, dropping SUMLEVEL & primary too, as they are of no use for further exploration.

```
In [15]: df_train.drop(columns=['BLOCKID', 'SUMLEVEL', 'primary'], inplace=True)
```

```
In [16]: df_test.drop(columns=['BLOCKID', 'SUMLEVEL', 'primary'], inplace=True)
```

```
In [17]: df_train.head(1)
```

```
Out[17]:
```

|        | COUNTYID | STATEID | state    | state_ab | city     | place    | type | zip_code | area_co |
|--------|----------|---------|----------|----------|----------|----------|------|----------|---------|
| UID    |          |         |          |          |          |          |      |          |         |
| 267822 | 53       | 36      | New York | NY       | Hamilton | Hamilton | City | 13346    | 3       |

1 rows × 76 columns

```
In [18]: df_test.head(1)
```

Out[18]:

| COUNTYID | STATEID | state | state_ab | city | place   | type                  | zip_code | area_ |
|----------|---------|-------|----------|------|---------|-----------------------|----------|-------|
| UID      |         |       |          |      |         |                       |          |       |
| 255504   | 163     | 26    | Michigan | MI   | Detroit | Dearborn Heights City | CDP      | 48239 |

1 rows x 76 columns

In [19]:

df\_train.info()

&lt;class 'pandas.core.frame.DataFrame'&gt;

Int64Index: 27321 entries, 267822 to 265371

Data columns (total 76 columns):

| #  | Column             | Non-Null | Count    | Dtype   |
|----|--------------------|----------|----------|---------|
| 0  | COUNTYID           | 27321    | non-null | int64   |
| 1  | STATEID            | 27321    | non-null | int64   |
| 2  | state              | 27321    | non-null | object  |
| 3  | state_ab           | 27321    | non-null | object  |
| 4  | city               | 27321    | non-null | object  |
| 5  | place              | 27321    | non-null | object  |
| 6  | type               | 27321    | non-null | object  |
| 7  | zip_code           | 27321    | non-null | int64   |
| 8  | area_code          | 27321    | non-null | int64   |
| 9  | lat                | 27321    | non-null | float64 |
| 10 | lng                | 27321    | non-null | float64 |
| 11 | ALand              | 27321    | non-null | float64 |
| 12 | AWater             | 27321    | non-null | int64   |
| 13 | pop                | 27321    | non-null | int64   |
| 14 | male_pop           | 27321    | non-null | int64   |
| 15 | female_pop         | 27321    | non-null | int64   |
| 16 | rent_mean          | 27007    | non-null | float64 |
| 17 | rent_median        | 27007    | non-null | float64 |
| 18 | rent_stdev         | 27007    | non-null | float64 |
| 19 | rent_sample_weight | 27007    | non-null | float64 |
| 20 | rent_samples       | 27007    | non-null | float64 |
| 21 | rent_gt_10         | 27007    | non-null | float64 |
| 22 | rent_gt_15         | 27007    | non-null | float64 |
| 23 | rent_gt_20         | 27007    | non-null | float64 |
| 24 | rent_gt_25         | 27007    | non-null | float64 |
| 25 | rent_gt_30         | 27007    | non-null | float64 |
| 26 | rent_gt_35         | 27007    | non-null | float64 |
| 27 | rent_gt_40         | 27007    | non-null | float64 |
| 28 | rent_gt_50         | 27007    | non-null | float64 |
| 29 | universe_samples   | 27321    | non-null | int64   |
| 30 | used_samples       | 27321    | non-null | int64   |
| 31 | hi_mean            | 27053    | non-null | float64 |
| 32 | hi_median          | 27053    | non-null | float64 |
| 33 | hi_stdev           | 27053    | non-null | float64 |
| 34 | hi_sample_weight   | 27053    | non-null | float64 |
| 35 | hi_samples         | 27053    | non-null | float64 |
| 36 | family_mean        | 27023    | non-null | float64 |
| 37 | family_median      | 27023    | non-null | float64 |

```

38 family_stdev                27023 non-null float64
39 family_sample_weight        27023 non-null float64
40 family_samples               27023 non-null float64
41 hc_mortgage_mean             26748 non-null float64
42 hc_mortgage_median           26748 non-null float64
43 hc_mortgage_stdev            26748 non-null float64
44 hc_mortgage_sample_weight    26748 non-null float64
45 hc_mortgage_samples          26748 non-null float64
46 hc_mean                      26721 non-null float64
47 hc_median                    26721 non-null float64
48 hc_stdev                     26721 non-null float64
49 hc_samples                   26721 non-null float64
50 hc_sample_weight             26721 non-null float64
51 home_equity_second_mortgage  26864 non-null float64
52 second_mortgage              26864 non-null float64
53 home_equity                  26864 non-null float64
54 debt                         26864 non-null float64
55 second_mortgage_cdf          26864 non-null float64
56 home_equity_cdf              26864 non-null float64
57 debt_cdf                     26864 non-null float64
58 hs_degree                    27131 non-null float64
59 hs_degree_male               27121 non-null float64
60 hs_degree_female             27098 non-null float64
61 male_age_mean                27132 non-null float64
62 male_age_median              27132 non-null float64
63 male_age_stdev               27132 non-null float64
64 male_age_sample_weight       27132 non-null float64
65 male_age_samples             27132 non-null float64
66 female_age_mean              27115 non-null float64
67 female_age_median            27115 non-null float64
68 female_age_stdev             27115 non-null float64
69 female_age_sample_weight     27115 non-null float64
70 female_age_samples           27115 non-null float64
71 pct_own                      27053 non-null float64
72 married                      27130 non-null float64
73 married_snp                  27130 non-null float64
74 separated                    27130 non-null float64
75 divorced                     27130 non-null float64
dtypes: float64(61), int64(10), object(5)
memory usage: 16.1+ MB

```

In [20]: `df_test.info()`

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 11709 entries, 255504 to 287763
Data columns (total 76 columns):
#   Column                Non-Null Count  Dtype
---  -
0   COUNTYID              11709 non-null  int64
1   STATEID               11709 non-null  int64
2   state                 11709 non-null  object
3   state_ab              11709 non-null  object
4   city                  11709 non-null  object
5   place                 11709 non-null  object
6   type                  11709 non-null  object
7   zip_code              11709 non-null  int64

```

|    |                             |       |          |         |
|----|-----------------------------|-------|----------|---------|
| 8  | area_code                   | 11709 | non-null | int64   |
| 9  | lat                         | 11709 | non-null | float64 |
| 10 | lng                         | 11709 | non-null | float64 |
| 11 | ALand                       | 11709 | non-null | int64   |
| 12 | AWater                      | 11709 | non-null | int64   |
| 13 | pop                         | 11709 | non-null | int64   |
| 14 | male_pop                    | 11709 | non-null | int64   |
| 15 | female_pop                  | 11709 | non-null | int64   |
| 16 | rent_mean                   | 11561 | non-null | float64 |
| 17 | rent_median                 | 11561 | non-null | float64 |
| 18 | rent_stdev                  | 11561 | non-null | float64 |
| 19 | rent_sample_weight          | 11561 | non-null | float64 |
| 20 | rent_samples                | 11561 | non-null | float64 |
| 21 | rent_gt_10                  | 11560 | non-null | float64 |
| 22 | rent_gt_15                  | 11560 | non-null | float64 |
| 23 | rent_gt_20                  | 11560 | non-null | float64 |
| 24 | rent_gt_25                  | 11560 | non-null | float64 |
| 25 | rent_gt_30                  | 11560 | non-null | float64 |
| 26 | rent_gt_35                  | 11560 | non-null | float64 |
| 27 | rent_gt_40                  | 11560 | non-null | float64 |
| 28 | rent_gt_50                  | 11560 | non-null | float64 |
| 29 | universe_samples            | 11709 | non-null | int64   |
| 30 | used_samples                | 11709 | non-null | int64   |
| 31 | hi_mean                     | 11587 | non-null | float64 |
| 32 | hi_median                   | 11587 | non-null | float64 |
| 33 | hi_stdev                    | 11587 | non-null | float64 |
| 34 | hi_sample_weight            | 11587 | non-null | float64 |
| 35 | hi_samples                  | 11587 | non-null | float64 |
| 36 | family_mean                 | 11573 | non-null | float64 |
| 37 | family_median               | 11573 | non-null | float64 |
| 38 | family_stdev                | 11573 | non-null | float64 |
| 39 | family_sample_weight        | 11573 | non-null | float64 |
| 40 | family_samples              | 11573 | non-null | float64 |
| 41 | hc_mortgage_mean            | 11441 | non-null | float64 |
| 42 | hc_mortgage_median          | 11441 | non-null | float64 |
| 43 | hc_mortgage_stdev           | 11441 | non-null | float64 |
| 44 | hc_mortgage_sample_weight   | 11441 | non-null | float64 |
| 45 | hc_mortgage_samples         | 11441 | non-null | float64 |
| 46 | hc_mean                     | 11419 | non-null | float64 |
| 47 | hc_median                   | 11419 | non-null | float64 |
| 48 | hc_stdev                    | 11419 | non-null | float64 |
| 49 | hc_samples                  | 11419 | non-null | float64 |
| 50 | hc_sample_weight            | 11419 | non-null | float64 |
| 51 | home_equity_second_mortgage | 11489 | non-null | float64 |
| 52 | second_mortgage             | 11489 | non-null | float64 |
| 53 | home_equity                 | 11489 | non-null | float64 |
| 54 | debt                        | 11489 | non-null | float64 |
| 55 | second_mortgage_cdf         | 11489 | non-null | float64 |
| 56 | home_equity_cdf             | 11489 | non-null | float64 |
| 57 | debt_cdf                    | 11489 | non-null | float64 |
| 58 | hs_degree                   | 11624 | non-null | float64 |
| 59 | hs_degree_male              | 11620 | non-null | float64 |
| 60 | hs_degree_female            | 11604 | non-null | float64 |
| 61 | male_age_mean               | 11625 | non-null | float64 |
| 62 | male_age_median             | 11625 | non-null | float64 |
| 63 | male_age_stdev              | 11625 | non-null | float64 |
| 64 | male_age_sample_weight      | 11625 | non-null | float64 |

```

65 male_age_samples          11625 non-null float64
66 female_age_mean          11613 non-null float64
67 female_age_median         11613 non-null float64
68 female_age_stdev          11613 non-null float64
69 female_age_sample_weight  11613 non-null float64
70 female_age_samples        11613 non-null float64
71 pct_own                   11587 non-null float64
72 married                   11625 non-null float64
73 married_snp               11625 non-null float64
74 separated                 11625 non-null float64
75 divorced                  11625 non-null float64
dtypes: float64(60), int64(11), object(5)
memory usage: 6.9+ MB

```

In [21]:

```

missing_train_cols = []

for col in df_train.columns:
    if df_train[col].isnull().sum() != 0:
        missing_train_cols.append(col)

print(missing_train_cols)

```

```

['rent_mean', 'rent_median', 'rent_stdev', 'rent_sample_weight', 'rent_samp
les', 'rent_gt_10', 'rent_gt_15', 'rent_gt_20', 'rent_gt_25', 'rent_gt_30',
'rent_gt_35', 'rent_gt_40', 'rent_gt_50', 'hi_mean', 'hi_median', 'hi_stdev
', 'hi_sample_weight', 'hi_samples', 'family_mean', 'family_median', 'famil
y_stdev', 'family_sample_weight', 'family_samples', 'hc_mortgage_mean', 'hc
_mortgage_median', 'hc_mortgage_stdev', 'hc_mortgage_sample_weight', 'hc_mo
rtgage_samples', 'hc_mean', 'hc_median', 'hc_stdev', 'hc_samples', 'hc_samp
le_weight', 'home_equity_second_mortgage', 'second_mortgage', 'home_equity'
, 'debt', 'second_mortgage_cdf', 'home_equity_cdf', 'debt_cdf', 'hs_degree'
, 'hs_degree_male', 'hs_degree_female', 'male_age_mean', 'male_age_median',
'male_age_stdev', 'male_age_sample_weight', 'male_age_samples', 'female_age
_mean', 'female_age_median', 'female_age_stdev', 'female_age_sample_weight'
, 'female_age_samples', 'pct_own', 'married', 'married_snp', 'separated', '
divorced']

```

In [22]:

```

missing_test_cols = []

for col in df_test.columns:
    if df_test[col].isnull().sum() != 0:
        missing_test_cols.append(col)

print(missing_test_cols)

```



```
['rent_mean', 'rent_median', 'rent_stdev', 'rent_sample_weight', 'rent_samples', 'rent_gt_10', 'rent_gt_15', 'rent_gt_20', 'rent_gt_25', 'rent_gt_30', 'rent_gt_35', 'rent_gt_40', 'rent_gt_50', 'hi_mean', 'hi_median', 'hi_stdev', 'hi_sample_weight', 'hi_samples', 'family_mean', 'family_median', 'family_stdev', 'family_sample_weight', 'family_samples', 'hc_mortgage_mean', 'hc_mortgage_median', 'hc_mortgage_stdev', 'hc_mortgage_sample_weight', 'hc_mortgage_samples', 'hc_mean', 'hc_median', 'hc_stdev', 'hc_samples', 'hc_sample_weight', 'home_equity_second_mortgage', 'second_mortgage', 'home_equity', 'debt', 'second_mortgage_cdf', 'home_equity_cdf', 'debt_cdf', 'hs_degree', 'hs_degree_male', 'hs_degree_female', 'male_age_mean', 'male_age_median', 'male_age_stdev', 'male_age_sample_weight', 'male_age_samples', 'female_age_mean', 'female_age_median', 'female_age_stdev', 'female_age_sample_weight', 'female_age_samples', 'pct_own', 'married', 'married_snp', 'separated', 'divorced']
```

Treating missing values by replacing it by mean as all the missing value col variables are numerical

```
In [23]: for col in df_train.columns:
         if col in (missing_train_cols):
             df_train[col].replace(np.nan, df_train[col].mean(), inplace=True)
```

```
In [24]: for col in df_test.columns:
         if col in (missing_test_cols):
             df_test[col].replace(np.nan, df_test[col].mean(), inplace=True)
```

```
In [25]: df_train.isnull().sum().any()
```

Out[25]: False

```
In [26]: df_test.isnull().sum().any()
```

Out[26]: False

## Exploratory Data Analysis (EDA):

Perform debt analysis. You may take the following steps:

- Explore the top 2,500 locations where the percentage of households with a second mortgage is the highest and percent ownership is above 10 percent.
- Visualize using geo-map.
- You may keep the upper limit for the percent of households with a second mortgage to 50 percent

```
In [27]: from pandasql import sqldf
         q1 = "select place,pct_own,second_mortgage,lat,lng from df_train where pct_
         pysqldf = lambda q: sqldf(q, globals())
         df_train_location_mort_pct=pysqldf(q1)
```

```
In [28]: df_train_location_mort_pct.head()
```

```
Out[28]:
```

|   | place           | pct_own | second_mortgage | lat       | lng        |
|---|-----------------|---------|-----------------|-----------|------------|
| 0 | Worcester City  | 0.20247 | 0.43363         | 42.254262 | -71.800347 |
| 1 | Harbor Hills    | 0.15618 | 0.31818         | 40.751809 | -73.853582 |
| 2 | Glen Burnie     | 0.22380 | 0.30212         | 39.127273 | -76.635265 |
| 3 | Egypt Lake-leto | 0.11618 | 0.28972         | 28.029063 | -82.495395 |
| 4 | Lincolnwood     | 0.14228 | 0.28899         | 41.967289 | -87.652434 |

```
In [29]: import plotly.express as px
         import plotly.graph_objects as go
```

In [106...


```
plt.figure(figsize=(15,30))
fig = go.Figure(data=go.Scattergeo(
    lat = df_train_location_mort_pct['lat'],
    lon = df_train_location_mort_pct['lng']),
)
fig.update_layout(
    geo=dict(
        scope = 'north america',
        showland = True,
        landcolor = "rgb(212, 212, 212)",
        subunitcolor = "rgb(255, 255, 255)",
        countrycolor = "rgb(255, 255, 255)",
        showlakes = True,
        lakecolor = "rgb(255, 255, 255)",
        showsubunits = True,
        showcountries = True,
        resolution = 50,
        projection = dict(
            type = 'conic conformal',
            rotation_lon = -100
        ),
        lonaxis = dict(
            showgrid = True,
            gridwidth = 0.5,
            range= [ -140.0, -55.0 ],
            dtick = 5
        ),
        lataxis = dict (
            showgrid = True,
            gridwidth = 0.5,
            range= [ 20.0, 60.0 ],
            dtick = 5
        )
    ),
    title='Top 2,500 locations with second mortgage is the highest and perc
fig.show()
```

<Figure size 1080x2160 with 0 Axes>

- Use the following bad debt equation:
  - $\text{Bad Debt} = P(\text{Second Mortgage} \cap \text{Home Equity Loan})$
  - $\text{Bad Debt} = \text{second\_mortgage} + \text{home\_equity} - \text{home\_equity\_second\_mortgage}$


```
In [31]: df_train["bad_debt"] = df_train['second_mortgage'] + df_train['home_equity']
```

```
In [32]: plt.figure(figsize=(15,10))
df_train['bins'] = pd.cut(df_train['bad_debt'],bins=[0,0.10,1], labels=["low","medium","high"])
df_train.groupby(['bins']).size().plot(kind='pie', subplots=True,startangle=0)
```



- Create pie charts to show overall debt and bad debt

```
In [33]: plt.figure(figsize=(15,10))
label=["Overall debt", "Bad Debt"]
values=[df_train["debt"].values.sum(), df_train["bad_debt"].values.sum()]
plt.pie(values, labels=label, autopct='%1.1f%%', explode=(0,.1), startangle=90)
plt.show()
```



- Create Box and whisker plot and analyze the distribution for 2nd mortgage, home equity, good debt, and bad debt for different cities

```
In [34]: df_train['good_debt']=df_train['debt']-df_train['bad_debt']
```

```
In [35]: df_train["city"].value_counts().sort_values(ascending=False)[:100]
```

```
Out[35]: Chicago      294
         Brooklyn    282
         Los Angeles  243
         Houston     222
         Philadelphia 165
         ...
         Norfolk     32
         Dayton      31
         Staten Island 31
         Manchester   31
         Franklin     31
         Name: city, Length: 100, dtype: int64
```


```
In [36]: plt.figure(figsize=(15,10))
import matplotlib.pyplot as plt
all_cities=df_train[['home_equity','second_mortgage','bad_debt', 'good_debt']]
all_cities.plot.box(grid=True)
plt.title('All Cities')
plt.show()
```

<Figure size 1080x720 with 0 Axes>



```
In [37]: plt.figure(figsize=(15,10))
Brooklyn=df_train[df_train['city']=='Brooklyn']
Brooklyn=Brooklyn[['home_equity','second_mortgage','bad_debt', 'good_debt']]
Brooklyn.plot.box(grid=True)
plt.title('Brooklyn')
plt.show()
```


&lt;Figure size 1080x720 with 0 Axes&gt;



In [38]:

```
plt.figure(figsize=(15,10))
Los_Angeles =df_train[df_train['city']=='Los Angeles']
Los_Angeles=Los_Angeles[['home_equity','second_mortgage','bad_debt', 'good_debt']]
Los_Angeles.plot.box(grid=True)
plt.title('Los Angeles')
plt.show()
```

&lt;Figure size 1080x720 with 0 Axes&gt;




In [39]:

```
plt.figure(figsize=(15,10))
Danville=df_train[df_train['city']=='Danville']
Danville=Danville[['home_equity','second_mortgage','bad_debt', 'good_debt']]
Danville.plot.box(grid=True)
plt.title('Danville')
plt.show()
```



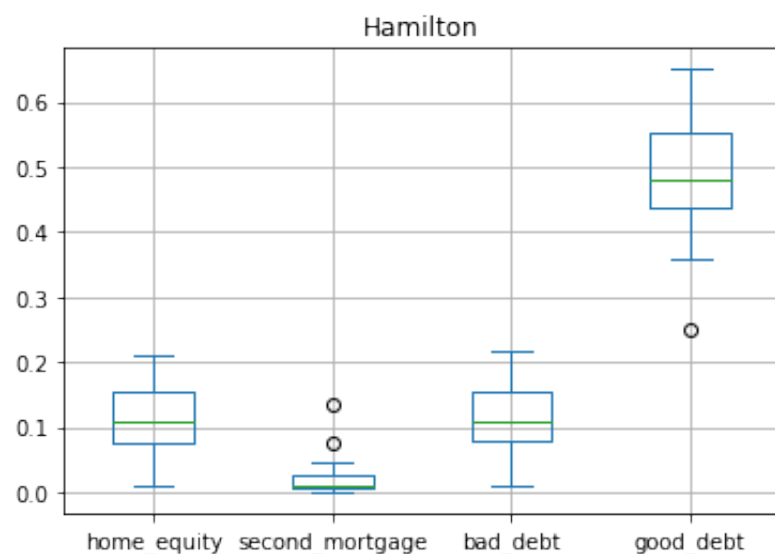
&lt;Figure size 1080x720 with 0 Axes&gt;



In [40]:

```
plt.figure(figsize=(15,10))
Hamilton=df_train[df_train['city']=='Hamilton']
Hamilton=Hamilton[['home_equity', 'second_mortgage', 'bad_debt', 'good_debt']]
Hamilton.plot.box(grid=True)
plt.title('Hamilton')
plt.show()
```


&lt;Figure size 1080x720 with 0 Axes&gt;



In [41]:

```
plt.figure(figsize=(15,10))
Houston =df_train[df_train['city']=='Houston']
Houston=Houston[['home_equity', 'second_mortgage', 'bad_debt', 'good_debt']]
Houston.plot.box(grid=True)
plt.title('Houston')
plt.show()
```


&lt;Figure size 1080x720 with 0 Axes&gt;



In [42]:

```
plt.figure(figsize=(15,10))
Manchester = df_train[df_train['city']=='Manchester']
Manchester=Manchester[['home_equity', 'second_mortgage', 'bad_debt', 'good_debt']]
Manchester.plot.box(grid=True)
plt.title('Manchester')
plt.show()
```

&lt;Figure size 1080x720 with 0 Axes&gt;



- Create a collated income distribution chart for family income, house hold income, and remaining income


In [43]:

```
plt.figure(figsize=(15,10))

plt.subplot(2,3,1)
sns.distplot(df_train['hi_mean'])
plt.title('Household Income')

plt.subplot(2,3,2)
sns.distplot(df_train['family_mean'])
plt.title('Family Income')

plt.subplot(2,3,3)
sns.distplot(df_train['family_mean']-df_train['hi_mean'])
plt.title('Remaining Income')
plt.show()
```



Perform EDA and come out with insights into population density and age. You may have to derive new fields (make sure to weight averages for accurate measurements):

- Use pop and ALand variables to create a new field called population density

In [44]:


```
df_train['pop_density']=df_train['pop']/df_train['ALand']
```

In [45]:

```
df_test['pop_density']=df_test['pop']/df_test['ALand']
```

In [46]:

```
plt.figure(figsize=(15,10))
sns.distplot(df_train['pop_density'])
plt.title('Population Density')
plt.show() # Very less density is noticed
```



- Use male\_age\_median, female\_age\_median, male\_pop, and female\_pop to create a new field called median age

```
In [47]: df_train['age_median']=(df_train['male_age_median']+df_train['female_age_med  
df_test['age_median']=(df_test['male_age_median']+df_test['female_age_media
```


- Visualize the findings using appropriate chart type

```
In [48]: df_train[['male_age_median','female_age_median','male_pop','female_pop','ac
```


```
Out[48]:
```

|        | male_age_median | female_age_median | male_pop | female_pop | age_median |
|--------|-----------------|-------------------|----------|------------|------------|
| UID    |                 |                   |          |            |            |
| 267822 | 44.00000        | 45.33333          | 2612     | 2618       | 44.666665  |
| 246444 | 32.00000        | 37.58333          | 1349     | 1284       | 34.791665  |
| 245683 | 40.83333        | 42.83333          | 3643     | 3238       | 41.833330  |
| 279653 | 48.91667        | 50.58333          | 1141     | 1559       | 49.750000  |
| 247218 | 22.41667        | 21.58333          | 2586     | 3051       | 22.000000  |

```
In [49]: plt.figure(figsize=(15,10))
sns.distplot(df_train['age_median']);
plt.title("Median Age")
plt.show()
```



```
In [50]: plt.figure(figsize=(15,10))
sns.boxplot(df_train['age_median']);
plt.title('Population Density')
plt.show()
```



Create bins for population into a new variable by selecting appropriate class interval so that the number of categories don't exceed 5 for the ease of analysis.

```
In [51]: df_train['pop_bins'] = pd.cut(df_train['pop'],bins=5,labels=['very_low', 'low', 'medium', 'high', 'very_high'])
```

```
In [52]: df_train['pop_bins'].value_counts()
```

```
Out[52]: very_low    27058
low           246
medium        9
high          7
very_high     1
Name: pop_bins, dtype: int64
```

- Analyze the married, separated, and divorced population for these population brackets


```
In [53]: df_train.groupby(by='pop_bins')[['married', 'separated', 'divorced']].count()
```

```
Out[53]:
```

|                  | married | separated | divorced |
|------------------|---------|-----------|----------|
| <b>pop_bins</b>  |         |           |          |
| <b>very_low</b>  | 27058   | 27058     | 27058    |
| <b>low</b>       | 246     | 246       | 246      |
| <b>medium</b>    | 9       | 9         | 9        |
| <b>high</b>      | 7       | 7         | 7        |
| <b>very_high</b> | 1       | 1         | 1        |

```
In [54]: plt.figure(figsize=(15,10))
df_train.groupby(by='pop_bins')[['married', 'separated', 'divorced']].count()
```

<Figure size 1080x720 with 0 Axes>



- Visualize using appropriate chart type

```
In [55]: df_train.groupby(by='pop_bins')[['married', 'separated', 'divorced']].agg()
```


```
Out[55]:
```

|                  | married  |          | separated |          | divorced |          |
|------------------|----------|----------|-----------|----------|----------|----------|
|                  | mean     | median   | mean      | median   | mean     | median   |
| <b>pop_bins</b>  |          |          |           |          |          |          |
| <b>very_low</b>  | 0.507548 | 0.524680 | 0.019126  | 0.013650 | 0.100504 | 0.096020 |
| <b>low</b>       | 0.584894 | 0.593135 | 0.015833  | 0.011195 | 0.075348 | 0.070045 |
| <b>medium</b>    | 0.655737 | 0.618710 | 0.005003  | 0.004120 | 0.065927 | 0.064890 |
| <b>high</b>      | 0.503359 | 0.335660 | 0.008141  | 0.002500 | 0.039030 | 0.010320 |
| <b>very_high</b> | 0.734740 | 0.734740 | 0.004050  | 0.004050 | 0.030360 | 0.030360 |

In [56]:

```
plt.figure(figsize=(15,10))  
df_train.groupby(by='pop_bins')[['married', 'separated', 'divorced']].agg()
```


&lt;Figure size 1080x720 with 0 Axes&gt;



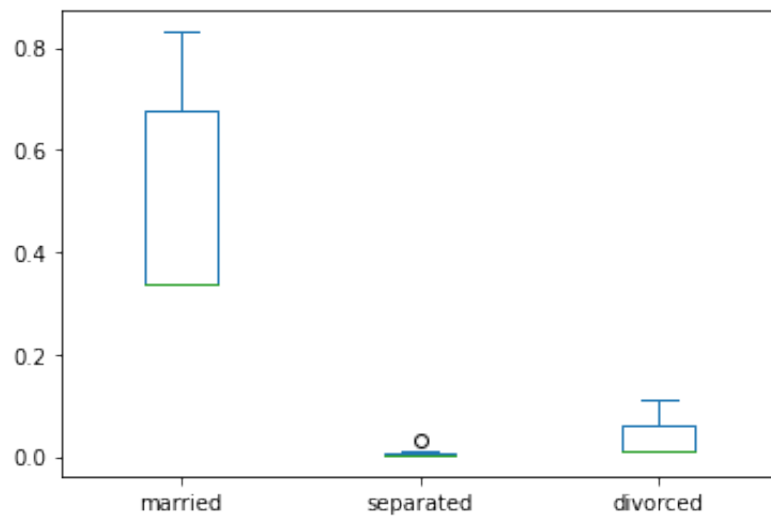
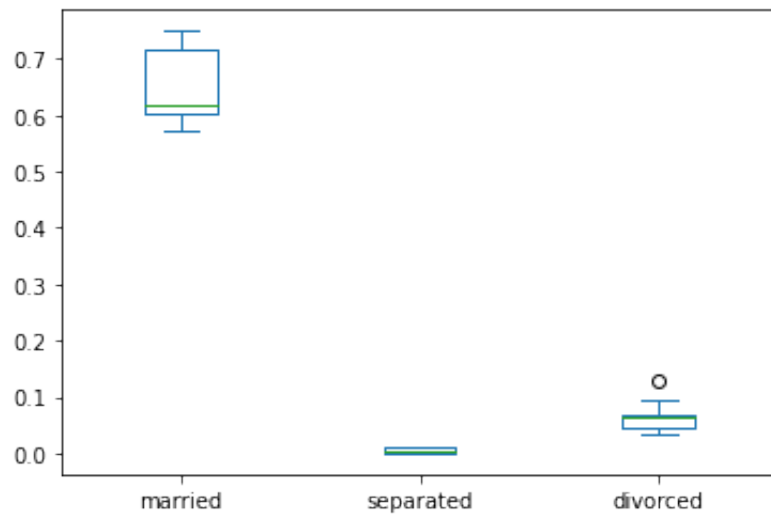
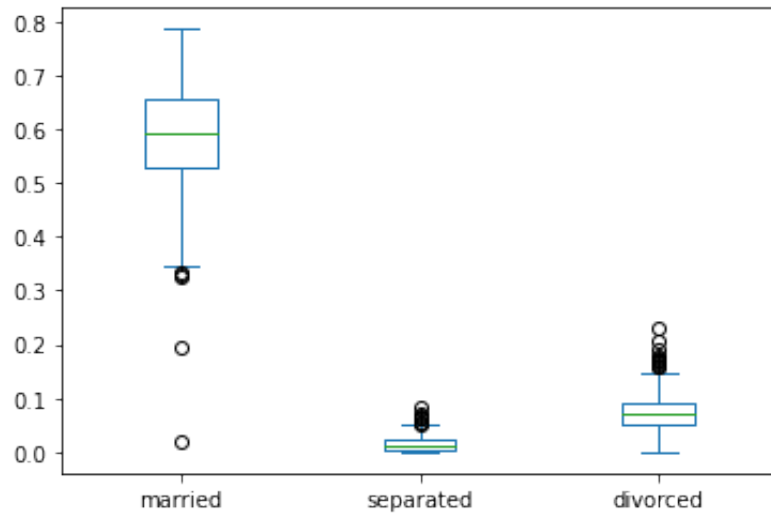
In [57]:


```
plt.figure(figsize=(15,10))  
df_train.groupby(by='pop_bins')[['married', 'separated', 'divorced']].plot
```

&lt;Figure size 1080x720 with 0 Axes&gt;










In [58]:

```
plt.figure(figsize=(10,5));
pop_bin_married = df_train.groupby(by='pop_bins')[['married', 'separated',
pop_bin_married.plot(figsize = (20,8));
plt.legend(loc='best');
plt.show();
```

&lt;Figure size 720x360 with 0 Axes&gt;



Please detail your observations for rent as a percentage of income at an overall level, and for different states.

In [59]:

```
rent_state_mean = df_train.groupby(by='state')['rent_mean'].agg(["mean"])
rent_state_mean.head()
```

Out [59]:

|            | mean        |
|------------|-------------|
| state      |             |
| Alabama    | 774.004927  |
| Alaska     | 1185.763570 |
| Arizona    | 1097.753511 |
| Arkansas   | 720.918575  |
| California | 1471.133857 |

| state      |             |
|------------|-------------|
| Alabama    | 774.004927  |
| Alaska     | 1185.763570 |
| Arizona    | 1097.753511 |
| Arkansas   | 720.918575  |
| California | 1471.133857 |

```
In [60]: income_state_mean = df_train.groupby(by='state')['family_mean'].agg(["mean",  
income_state_mean.head()
```

Out [60]:

|            | mean         |
|------------|--------------|
| state      |              |
| Alabama    | 67030.064213 |
| Alaska     | 92136.545109 |
| Arizona    | 73328.238798 |
| Arkansas   | 64765.377850 |
| California | 87655.470820 |

| state      |              |
|------------|--------------|
| Alabama    | 67030.064213 |
| Alaska     | 92136.545109 |
| Arizona    | 73328.238798 |
| Arkansas   | 64765.377850 |
| California | 87655.470820 |

```
In [61]: rent_per_of_income = rent_state_mean['mean']/income_state_mean['mean']
```

```
In [62]: rent_per_of_income
```

```
Out[62]: state
Alabama      0.011547
Alaska       0.012870
Arizona      0.014970
Arkansas     0.011131
California   0.016783
Colorado     0.013529
Connecticut  0.012637
Delaware     0.012929
District of Columbia 0.013198
Florida      0.015772
Georgia      0.013161
Hawaii       0.018224
Idaho        0.011957
Illinois     0.012620
Indiana      0.012022
Iowa         0.009940
Kansas       0.011066
Kentucky     0.011068
Louisiana    0.012160
Maine        0.011674
Maryland     0.013947
Massachusetts 0.012312
Michigan     0.012766
Minnesota    0.011058
Mississippi  0.012428
Missouri     0.011670
Montana      0.010789
Nebraska     0.010912
Nevada       0.015242
New Hampshire 0.011949
New Jersey   0.013678
New Mexico   0.012330
New York     0.014410
North Carolina 0.012166
North Dakota 0.009303
Ohio         0.011401
Oklahoma     0.011632
Oregon       0.013253
Pennsylvania 0.011902
Puerto Rico 0.015133
Rhode Island 0.012292
South Carolina 0.012657
South Dakota 0.009192
Tennessee    0.012286
Texas        0.012899
Utah         0.013192
Vermont      0.011743
Virginia     0.014050
Washington   0.013352
West Virginia 0.010341
Wisconsin    0.011189
Wyoming      0.010785
Name: mean, dtype: float64
```

```
In [63]: sum(df_train['rent_mean'])/sum(df_train['family_mean'])
```

```
Out[63]: 0.013358170721473864
```


Perform correlation analysis for all the relevant variables by creating a heatmap. Describe your findings.

```
In [64]: df_train[['hc_mortgage_mean', 'ALand', 'pop', 'rent_mean', 'hi_mean', 'hc_mean',
```

```
Out[64]:
```

|                  | hc_mortgage_mean | ALand     | pop       | rent_mean | hi_mean   | hc_mean |
|------------------|------------------|-----------|-----------|-----------|-----------|---------|
| hc_mortgage_mean | 1.000000         | -0.056334 | 0.110659  | 0.750081  | 0.763128  | 0.7     |
| ALand            | -0.056334        | 1.000000  | -0.033743 | -0.067169 | -0.028435 | -0.0    |
| pop              | 0.110659         | -0.033743 | 1.000000  | 0.160590  | 0.166913  | 0.0     |
| rent_mean        | 0.750081         | -0.067169 | 0.160590  | 1.000000  | 0.753920  | 0.5     |
| hi_mean          | 0.763128         | -0.028435 | 0.166913  | 0.753920  | 1.000000  | 0.6     |
| hc_mean          | 0.795012         | -0.056723 | 0.051515  | 0.594499  | 0.675090  | 1.0     |
| family_mean      | 0.759805         | -0.027897 | 0.128173  | 0.701019  | 0.960624  | 0.6     |
| hs_degree        | 0.333336         | -0.002293 | 0.049238  | 0.362944  | 0.580284  | 0.3     |
| debt             | 0.390902         | -0.115591 | 0.231013  | 0.432481  | 0.418408  | 0.2     |
| home_equity      | 0.466481         | -0.079494 | 0.099352  | 0.408837  | 0.469863  | 0.3     |

```
In [65]: plt.figure(figsize=(17,7))
sns.heatmap(df_train[['hc_mortgage_mean', 'ALand', 'pop', 'rent_mean', 'hi_mean',
plt.show()
```



'rent\_mean', 'hi\_mean', 'hc\_mean', 'family\_mean' has a good corr with our target variable - hc\_mortgage\_mean

## Project Task: Week 2

### Data Pre-processing:

The economic multivariate data has a significant number of measured variables. The goal is to find where the measured variables depend on a number of smaller unobserved common factors or latent variables.

Each variable is assumed to be dependent upon a linear combination of the common factors, and the coefficients are known as loadings. Each measured variable also includes a component due to independent random variability, known as "specific variance" because it is specific to one variable. Obtain the common factors and then plot the loadings. Use factor analysis to find latent variables in our dataset and gain insight into the linear relationships in the data.

Following are the list of latent variables:

Highschool graduation rates

Median population age

Second mortgage statistics

Percent own

Bad debt expense

```
In [66]: from sklearn.decomposition import FactorAnalysis  
        from factor_analyzer import FactorAnalyzer
```

```
In [67]: df_train.describe()
```

Out [67]:

|       | COUNTYID     | STATEID      | zip_code     | area_code    | lat          |              |
|-------|--------------|--------------|--------------|--------------|--------------|--------------|
| count | 27321.000000 | 27321.000000 | 27321.000000 | 27321.000000 | 27321.000000 | 27321.000000 |
| mean  | 85.646426    | 28.271806    | 50081.999524 | 596.507668   | 37.508813    | -91.2883     |
| std   | 98.333097    | 16.392846    | 29558.115660 | 232.497482   | 5.588268     | 16.3438      |
| min   | 1.000000     | 1.000000     | 602.000000   | 201.000000   | 17.929085    | -165.4538    |
| 25%   | 29.000000    | 13.000000    | 26554.000000 | 405.000000   | 33.899064    | -97.8160     |
| 50%   | 63.000000    | 28.000000    | 47715.000000 | 614.000000   | 38.755183    | -86.5543     |
| 75%   | 109.000000   | 42.000000    | 77093.000000 | 801.000000   | 41.380606    | -79.7825     |
| max   | 840.000000   | 72.000000    | 99925.000000 | 989.000000   | 67.074017    | -65.3793     |

8 rows x 75 columns

In [68]:

```
fa = FactorAnalyzer(n_factors=5)
```

In [69]:

```
fa.fit_transform(df_train.select_dtypes(exclude=('object', 'category')))
```

Out [69]:

```
array([[ -0.39933034,  0.55583772,  1.07093896, -1.08698579,  0.65355193],
       [ -0.99248908, -0.57075893, -0.12202251,  0.10554693,  0.28386921],
       [  0.02533901,  1.21506669,  0.4946933 , -0.50562619, -0.28848015],
       ...,
       [  0.02046437, -0.70649548,  0.81319778, -1.37996186,  0.00873314],
       [  2.51731673,  3.10777987,  1.14759443, -0.0630024 , -1.63083959],
       [ -0.33101021, -0.23542995, -1.63171941,  0.17824411, -0.12268557]])
```

In [70]:

```
fa.loadings_
```

Out [70]:

```
array([[ -0.11482487,  0.01936373, -0.0245545 , -0.06169006,  0.03812136],
       [ -0.11040899,  0.01429921,  0.02466482, -0.14796223,  0.11258263],
       [ -0.0891633 ,  0.04864149, -0.12749732, -0.04931909, -0.11805326],
       [  0.01614823,  0.0188282 ,  0.00574863,  0.02659233, -0.00973247],
       [  0.09063446, -0.09926159, -0.05333377, -0.13305276, -0.14641972],
       [ -0.00541231, -0.03861633,  0.13839861,  0.00876519,  0.1216676 ],
       [ -0.04186143, -0.02024517,  0.03644117, -0.09300196,  0.06439296],
       [ -0.00198676, -0.0150142 , -0.00250851, -0.04444931,  0.02563788],
       [  0.07644546,  0.95538373, -0.08288847, -0.00717471, -0.05369428],
       [  0.07116302,  0.91659973, -0.10342851, -0.02796311, -0.05283109],
       [  0.07805518,  0.94595683, -0.05801542,  0.014232 , -0.05251385],
       [  0.76054105,  0.00785577, -0.03725023,  0.11387679, -0.14404771],
       [  0.70885225,  0.00394816, -0.04466907,  0.10817145, -0.15591822],
       [  0.70643093,  0.02688275, -0.02517887,  0.10292387,  0.06883075],
       [ -0.121171 ,  0.34301813, -0.51395397, -0.04394071,  0.31530346],
       [  0.24323025,  0.44469406, -0.67418609, -0.02865328,  0.33289811],
       [ -0.04675332,  0.03348746,  0.03198187,  0.44044074, -0.1698257 ],
       [ -0.02516145,  0.01631269,  0.04136891,  0.66998302, -0.16155432],
       [ -0.03634064, -0.01567031,  0.06956415,  0.82915629, -0.09407074],
       [ -0.04712842, -0.03348549,  0.09423677,  0.91747318, -0.04414534],
       [ -0.05620306, -0.04149989,  0.1177372 ,  0.94558341, -0.02041133],
```

```
[ -0.0412783 , -0.04959822, 0.12201166, 0.92570715, 0.00189047 ],
[ -0.03738612, -0.05602061, 0.11023791, 0.88082495, 0.01283847 ],
[ -0.02031965, -0.06992452, 0.07639516, 0.77345015, 0.03171296 ],
[ 0.2251697 , 0.47338185, -0.64943552, -0.02812309, 0.35439151 ],
[ 0.24516516, 0.45397088, -0.66203053, -0.03064164, 0.32930185 ],
[ 0.77424469, 0.0458412 , 0.15426491, -0.20366713, -0.16645681 ],
[ 0.69788075, 0.04545261, 0.14737086, -0.21735663, -0.22413606 ],
[ 0.85744941, 0.04437585, 0.15809972, -0.11959275, 0.02692194 ],
[ -0.21235552, 0.85071284, -0.06515656, 0.06559846, 0.23273441 ],
[ 0.14775965, 0.95466245, 0.01314345, -0.04816854, 0.09715259 ],
[ 0.82240955, 0.03269749, 0.16390804, -0.20341937, -0.08254877 ],
[ 0.78596814, 0.02645931, 0.15603318, -0.20627097, -0.09951158 ],
[ 0.80935346, 0.04480529, 0.13331803, -0.1080682 , 0.05512187 ],
[ -0.33493419, 0.86485945, 0.03359967, 0.08892098, 0.04245147 ],
[ 0.04675907, 0.9327849 , 0.15952872, -0.02606483, -0.10045735 ],
[ 0.97737122, -0.03058639, -0.12023358, 0.04422083, 0.06391006 ],
[ 0.95753903, -0.03670088, -0.13489606, 0.04450991, 0.05518235 ],
[ 0.81891124, 0.0062711 , 0.0606408 , 0.01965969, 0.13413668 ],
[ -0.42447833, 0.71211816, 0.36471085, -0.07278991, -0.28356804 ],
[ 0.06479058, 0.71798809, 0.30113511, -0.04925706, -0.36533169 ],
[ 0.90925698, -0.05005301, -0.06873 , -0.00156445, 0.15344429 ],
[ 0.8716504 , -0.0497007 , -0.07950161, -0.00230469, 0.141696 ],
[ 0.75826576, 0.00185445, 0.02736492, 0.00374285, 0.25783346 ],
[ -0.11345437, 0.61298275, 0.62349709, -0.01999399, 0.284585 ],
[ -0.33349554, 0.5640579 , 0.58243645, -0.02351225, 0.25401294 ],
[ -0.13962631, -0.02038193, -0.099815 , 0.10703584, -0.61108334 ],
[ -0.1155693 , -0.02672498, -0.10052413, 0.12267069, -0.6194846 ],
[ 0.26332311, -0.02989007, 0.02773139, 0.09313835, -0.58884017 ],
[ 0.17605066, 0.06500774, -0.27623933, 0.01755758, -0.70238435 ],
[ 0.09571295, -0.05641886, -0.0878796 , -0.09217904, 0.64282047 ],
[ -0.27816504, -0.00108849, -0.08724621, -0.09099697, 0.60287301 ],
[ -0.18822894, -0.05993517, 0.3221272 , -0.01484872, 0.70805651 ],
[ 0.39096931, 0.05805001, 0.26868167, -0.2206138 , -0.17910889 ],
[ 0.40449327, 0.06000692, 0.23582956, -0.21037746, -0.16855004 ],
[ 0.35015101, 0.05083956, 0.28323451, -0.21703079, -0.17351309 ],
[ 0.24845274, -0.04071472, 0.7828657 , 0.09032636, 0.37052567 ],
[ 0.25210151, -0.02708951, 0.80625444, 0.07208958, 0.28641364 ],
[ -0.05920176, 0.06895631, 0.57728575, 0.08529851, 0.12110191 ],
[ 0.05379624, 0.81586993, -0.1714922 , -0.0147333 , -0.04544873 ],
[ 0.06864777, 0.92118309, -0.10138082, -0.02787599, -0.05733179 ],
[ 0.20884889, -0.03872551, 0.76992343, 0.13951893, 0.37840595 ],
[ 0.20114043, -0.02576923, 0.83100793, 0.12749903, 0.29865452 ],
[ -0.09355638, 0.06379344, 0.46282287, 0.07177015, 0.13940071 ],
[ 0.05905288, 0.87591732, -0.1452326 , 0.02192923, -0.05554446 ],
[ 0.07559661, 0.95252905, -0.0559462 , 0.01535826, -0.05728285 ],
[ -0.04117527, 0.10599545, 0.80950398, -0.04276026, -0.26442385 ],
[ 0.17123066, 0.18784317, 0.57209879, -0.12048879, -0.12264931 ],
[ -0.06001846, -0.06689816, -0.28697978, 0.12758949, 0.1806349 ],
[ -0.15149224, -0.06790769, -0.16008983, 0.12294432, 0.14495427 ],
[ -0.34836642, -0.05022182, 0.13810517, 0.02668981, 0.12886669 ],
[ 0.26104601, -0.03305257, 0.02283549, 0.10263406, -0.60019879 ],
[ 0.07779707, 0.09682759, -0.32690032, -0.03383985, -0.46819902 ],
[ 0.35337036, -0.00539003, -0.42547832, 0.05776243, 0.27393594 ],
[ 0.23903134, -0.02686691, 0.86480585, 0.10861162, 0.31067539 ] ]
```



## Data Modeling :

Build a linear Regression model to predict the total monthly expenditure for home mortgages loan.

Please refer deplotment\_RE.xlsx. Column hc\_mortgage\_mean is predicted variable. This is the mean monthly mortgage and owner costs of specified geographical location.

Note: Exclude loans from prediction model which have NaN (Not a Number) values for hc\_mortgage\_mean.

a) Run a model at a Nation level. If the accuracy levels and R square are not satisfactory proceed to below step.

b) Run another model at State level. There are 52 states in USA.

c) Keep below considerations while building a linear regression model:

Variables should have significant impact on predicting Monthly mortgage and owner costs

Utilize all predictor variable to start with initial hypothesis

R square of 60 percent and above should be achieved

Ensure Multi-collinearity does not exist in dependent variables

Test if predicted variable is normally distributed

```
In [71]: df_train.columns
```

```
Out[71]: Index(['COUNTYID', 'STATEID', 'state', 'state_ab', 'city', 'place', 'type',
              'zip_code', 'area_code', 'lat', 'lng', 'ALand', 'AWater', 'pop',
              'male_pop', 'female_pop', 'rent_mean', 'rent_median', 'rent_stdev',
              'rent_sample_weight', 'rent_samples', 'rent_gt_10', 'rent_gt_15',
              'rent_gt_20', 'rent_gt_25', 'rent_gt_30', 'rent_gt_35', 'rent_gt_40',
              ,
              'rent_gt_50', 'universe_samples', 'used_samples', 'hi_mean',
              'hi_median', 'hi_stdev', 'hi_sample_weight', 'hi_samples',
              'family_mean', 'family_median', 'family_stdev', 'family_sample_weight',
              ,
              'family_samples', 'hc_mortgage_mean', 'hc_mortgage_median',
              'hc_mortgage_stdev', 'hc_mortgage_sample_weight', 'hc_mortgage_samples',
              ,
              'hc_mean', 'hc_median', 'hc_stdev', 'hc_samples', 'hc_sample_weight',
              ,
              'home_equity_second_mortgage', 'second_mortgage', 'home_equity', 'debt',
              ,
              'second_mortgage_cdf', 'home_equity_cdf', 'debt_cdf', 'hs_degree',
              'hs_degree_male', 'hs_degree_female', 'male_age_mean',
              'male_age_median', 'male_age_stdev', 'male_age_sample_weight',
              'male_age_samples', 'female_age_mean', 'female_age_median',
              'female_age_stdev', 'female_age_sample_weight', 'female_age_samples',
              ,
              'pct_own', 'married', 'married_snp', 'separated', 'divorced',
              'bad_debt', 'bins', 'good_debt', 'pop_density', 'age_median',
              'pop_bins'],
              dtype='object')
```

```
In [72]: df_train['type'].unique()
```

```
Out[72]: array(['City', 'Urban', 'Town', 'CDP', 'Village', 'Borough'], dtype=object)
```

```
In [73]: type_dict = {'type':{'City':1, 'Urban':2, 'Town':3, 'CDP':4, 'Village':5, 'Borough':6}}
```

```
In [74]: df_train.replace(type_dict, inplace=True)
```

```
In [75]: df_test.replace(type_dict, inplace=True)
```

```
In [76]: df_train['type'].unique()
```

```
Out[76]: array([1, 2, 3, 4, 5, 6])
```

```
In [77]: feature_cols = ['COUNTYID', 'STATEID', 'zip_code', 'type', 'pop', 'family_mean',
                        'hs_degree', 'age_median', 'pct_own', 'married', 'separated', 'divorced', 'bad_debt', 'good_debt', 'pop_density', 'pop_bins']
```

```
In [78]: x_train = df_train[feature_cols]
        y_train = df_train['hc_mortgage_mean']
```

```
In [79]: x_test = df_test[feature_cols]
        y_test = df_test['hc_mortgage_mean']
```

```
In [80]: x_train.shape, y_train.shape
```

```
Out[80]: ((27321, 15), (27321,))
```

```
In [81]: x_test.shape, y_test.shape
```

```
Out[81]: ((11709, 15), (11709,))
```

```
In [82]: from sklearn.preprocessing import StandardScaler
        from sklearn.linear_model import LinearRegression
        from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error
```

```
In [83]: x_test.head()
```

```
Out[83]:
```

|        | COUNTYID | STATEID | zip_code | type | pop  | family_mean  | second_mortgage | hon |
|--------|----------|---------|----------|------|------|--------------|-----------------|-----|
| UID    |          |         |          |      |      |              |                 |     |
| 255504 | 163      | 26      | 48239    | 4    | 3417 | 53802.87122  | 0.06443         |     |
| 252676 | 1        | 23      | 4210     | 1    | 3796 | 85642.22095  | 0.01175         |     |
| 276314 | 15       | 42      | 14871    | 6    | 3944 | 65694.06582  | 0.01316         |     |
| 248614 | 231      | 21      | 42633    | 1    | 2508 | 44156.38709  | 0.00995         |     |
| 286865 | 355      | 48      | 78410    | 3    | 6230 | 123527.02420 | 0.00000         |     |

```
In [84]: SC = StandardScaler()
```

```
In [85]: x_train_Scaled = SC.fit_transform(X_train)
        x_test_Scaled = SC.fit_transform(X_test)
```

```
In [86]: Lr = LinearRegression()
```

```
In [87]: Lr.fit(X_train_Scaled,y_train)
```

```
Out[87]: LinearRegression()
```

```
In [88]: y_pred = Lr.predict(X_test_Scaled)
```

```
In [89]: r2_score(y_test,y_pred) #R2 square
```

```
Out[89]: 0.7348210754610929
```

```
In [90]: np.sqrt(mean_squared_error(y_test,y_pred)) #RMSE
```

```
Out[90]: 323.1018894984635
```

**b) Run another model at State level. There are 52 states in USA.**

```
In [91]: state= df_train['STATEID'].nunique()
```

```
In [92]: state= df_train['STATEID'].unique()  
state
```

```
Out[92]: array([36, 18, 72, 20,  1, 48, 45,  6,  5, 24, 17, 19, 47, 32, 22,  8, 44,  
          28, 34, 41,  4, 12, 55, 42, 37, 51, 26, 39, 40, 13, 16, 46, 27, 29,  
          53, 56,  9, 54, 21, 25, 11, 15, 30,  2, 33, 49, 50, 31, 38, 35, 23,  
          10])
```

In [93]:

```
for i in [11,33,35]:
    print("State_ID:",i)

    X_train_nation = df_train[df_train['COUNTYID']==i][feature_cols]
    y_train_nation = df_train[df_train['COUNTYID']==i]['hc_mortgage_mean']

    X_test_nation = df_test[df_test['COUNTYID']==i][feature_cols]
    y_test_nation = df_test[df_test['COUNTYID']==i]['hc_mortgage_mean']

    X_train_Scaled_nation = SC.fit_transform(X_train_nation)
    X_test_Scaled_nation = SC.fit_transform(X_test_nation)

    Lr.fit(X_train_Scaled_nation, y_train_nation)

    y_pred_nation = Lr.predict(X_test_Scaled_nation)

    print("Overall R2 score of linear regression model for state," ,i," :-"
    print("Overall RMSE of linear regression model for state," ,i," :-" ,np.s
    print("\n")
```

State\_ID: 11

Overall R2 score of linear regression model for state, 11 :- 0.746485716944  
4445

Overall RMSE of linear regression model for state, 11 :- 238.10563068257605

State\_ID: 33

Overall R2 score of linear regression model for state, 33 :- 0.861561420773  
1607

Overall RMSE of linear regression model for state, 33 :- 211.13273527746531


State\_ID: 35

Overall R2 score of linear regression model for state, 35 :- 0.722243579050  
943


Overall RMSE of linear regression model for state, 35 :- 255.6042328589991

In [94]:

```
sns.distplot(y_pred)
plt.show()
```




```
In [95]: sns.distplot(y_pred_nation)
plt.show()
```



```
In [96]: residuals=y_test-y_pred
residuals
```

```
Out[96]: UID
255504    281.969088
252676   -69.935775
276314    190.761969
248614  -157.290627
286865   -9.887017
...
238088   -67.541646
242811   -41.578757
250127  -127.427569
241096  -330.820475
287763    217.760642
Name: hc_mortgage_mean, Length: 11709, dtype: float64
```


```
In [97]: sns.distplot(residuals)
plt.show()
```




```
In [98]: residuals_nation=y_test_nation-y_pred_nation
residuals_nation
```

```
Out[98]: UID
271383    -71.037462
271503    -92.539838
288357     84.266127
264513    -66.054827
288371    456.709228
...
288512     44.115107
280441    -48.716870
288475    -49.259951
253291     74.474223
288425    201.479714
Name: hc_mortgage_mean, Length: 187, dtype: float64
```


```
In [99]: sns.distplot(residuals_nation)
plt.show()
```



```
In [100... plt.scatter(residuals,y_pred);
```



```
In [101... plt.scatter(residuals_nation,y_pred_nation);
```





## Data Reporting:

- Create a dashboard in tableau by choosing appropriate chart types and metrics useful for the business. The dashboard must entail the following:
- Box plot of distribution of average rent by type of place (village, urban, town, etc.).
- Pie charts to show overall debt and bad debt.
- Explore the top 2,500 locations where the percentage of households with a second mortgage is the highest and percent ownership is above 10 percent. Visualize using geo-map.
- Heat map for correlation matrix.
- Pie chart to show the population distribution across different types of places (village, urban, town etc.).

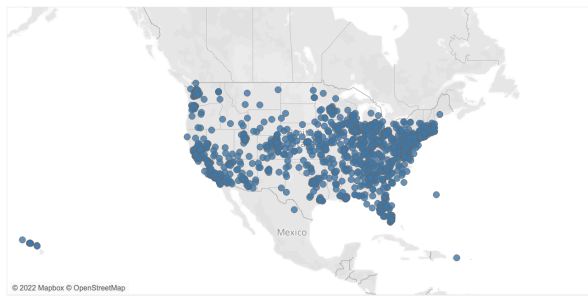
```
In [102... df_train_location_mort_pct.to_excel('df_train_location_mort_pct.xlsx')
```

```
In [103... df_train.to_excel('df_train_.xlsx')
```


```
In [105... df_test.to_excel('df_test_.xlsx')
```

```
In [107... df_train.bad_debt.to_excel('bad_debt.xlsx')
```


Explore the top 2,500 locations where the percentage of households with a second mortgage is the highest and percent ownership is above 10 percent. Visualize using geo-map.




Box plot of distribution of average rent by type of place (village, urban, town, etc.).




Heat map for correlation matrix.(Hi Mean, Hc Mean)




Heat map for correlation matrix.(Family Mean, Rent Mean)



Pie chart to show the population distribution across different types of places (village, urban, town etc.).



Pie charts to show overall debt and bad debt.



tribution of averag... Pie charts to show overall debt a... Explore the top 2,500 locations ... Heat map for correlation matrix... Heat map for correlation matrix... Pie chart to show the population... Dashboard 2

<https://public.tableau.com/app/profile/rushikesh.khankar/viz/RealEStatisticalAnalysisCapProject/Dashboard2>

Thank You